

• THE PROBLEM

Email QA is still mostly manual.

Marketing teams repeatedly test the same fragile journey logic before every launch: inbox checks, branch checks, CTA checks, tracking checks, personalization checks, and compliance checks.

"The challenge is proving the whole customer journey behaves correctly for every persona."

1

Set up seed contacts

Create aliases/personas for clickers, non-clickers, product preferences, and edge cases.

repeat

2

Wait and refresh inboxes

Check whether each message arrives at the right time for the right contact.

slow

3

Click through paths

Manually test CTAs without accidentally triggering the wrong branch.

fragile

4

Write the QA result

Summarize screenshots, email subjects, broken links, and launch blockers.

manual

One campaign creates dozens of tiny checks.

MANUAL QA

TASK	HUMAN ACTION
Branch logic	Compare expected vs actual emails for every persona.
CTA checks	Open links, avoid tracking side effects, record result.
Personalization	Scan copy for raw tokens and missing fallbacks.
Compliance	Find unsubscribe link and verify destination.
Report	Manually turn observations into launch guidance.

INBOXFLOW AGENT

TASK	AGENT ACTION
Branch logic	Builds expected persona paths and compares observed delivery.
CTA checks	Clicks only when persona behavior requires it.
Personalization	Flags unresolved tokens with email proof.
Compliance	Checks unsubscribe presence and destination intent.
Report	Generates launch readiness, replay, fixes, and evidence.

"InboxFlow removes repetitive observation work so humans can focus on fixing the journey."

The expensive failures are invisible until they ship.

Wrong branch

A non-clicker receives a click-only follow-up.

This means the actual customer journey does not match the intended automation logic.

Broken revenue path

The primary CTA lands on the wrong page or a dead destination.

The email can look fine while the conversion path is unusable.

Brand and compliance risk

Raw tokens, missing tracking, or broken unsubscribe flows go live.

These are small QA misses that become very visible to customers.

An autonomous launch gate for lifecycle campaigns.

Parse

Turns plain English into personas, branches, and timed steps.

Watch

Monitors the seed inbox for each persona alias.

Act

Clicks, opens, or holds back based on persona behavior.

Probe

Checks landing pages while avoiding false click tracking.

Reason

Uses semantic checks for copy, intent, and page alignment.

Report


Writes a launch-readiness decision with evidence.

"Gemini interprets intent. The backend executes safely. The report is evidence-backed."

• DEMO REVEAL

The agent catches the journey, not just the email.

The strongest demo moment is the mismatch between what the marketer expected and what the customer actually received.



Expected result

The non-clicker should receive Reminder 2B, not the click-only follow-up.

From manual QA labor to actionable fix list.

Less

Inbox refreshing, screenshot collection, and spreadsheet tracking.

More

Time spent fixing the automation, copy, links, and compliance issues.

Clear

One launch decision with owners: Journey builder, Content, Tracking, Web/Dev, Compliance.

"Do not launch until branch logic, CTA destination, and personalization issues are fixed."

Some QA cannot be solved by status codes.

A page can return successfully and still be the wrong page. A subject can be valid text and still mislead the recipient. That is where semantic reasoning helps.

Deterministic checks

- ✓ Did the expected email arrive?
- ✓ Did a persona receive the wrong branch?
- ✓ Are UTM parameters present?
- ✓ Are raw personalization tokens visible?

Gemini reasoning

- ✓ Does the CTA destination match the campaign promise?
- ✓ Does the unsubscribe destination actually look like opt-out?
- ✓ Do subject, body, CTA, and goal align?

The browser is the cockpit. Vultr runs the agent.

For the hackathon minimum, Vultr Cloud Compute gives us the public, persistent runtime for the server-side agent loop.

Public demo URL

Judges can open one app URL instead of watching a local laptop-only demo.

Vultr Cloud Compute

Runs Express API, step runner, Gmail sync, semantic report generation, and persisted run state.

External services

Gemini interprets and reasons. Gmail supplies the seed inbox evidence.

"Vultr is our stateful agent runtime. It hosts the backend that waits, syncs, acts, validates, and records the launch decision."

Why this is more than another AI wrapper.

Enterprise utility

Marketing operations teams already do this manually before launch. The pain is real and recurring.

Agentic workflow

The agent perceives inbox state, acts as personas, waits for timing windows, validates outcomes, and writes a report.

Measurable value

Fewer manual test passes, fewer missed campaign errors, faster launch confidence, clearer owner-based fixes.

"InboxFlow validates whether an automated customer journey is safe to launch."

The demo is a simple before-and-after story.

- 1 Open InboxFlow**
Public app opens from the Vultr-hosted deployment.

- 2 Start a demo campaign**
A plain-language journey becomes personas, branches, and timed steps.

- 3 Show the activity timeline**
The agent runs the repetitive QA loop: sync, wait, act, sync again, validate.

- 4 Open the report**
The output is a launch decision with evidence, not a generated opinion.

- 5 Reveal the blocker**
The non-clicker received the clicker follow-up, so the campaign should not launch.

• CLOSING LINE

InboxFlow makes email launch QA auditable.

It removes repetitive human checking, catches customer-facing journey failures, and gives teams a clear launch decision.

Before

Manual inbox checking, spreadsheet comparison, screenshots, subjective launch confidence.

After

Autonomous test run, persona replay, evidence pack, owner-based fixes, deterministic launch verdict.

"Gemini interprets intent. Vultr runs the agent. InboxFlow blocks unsafe launches."